

DAMG 6210

Database Design

Nik Bear Brown

@NikBearBrown

SQL Constraints & Triggers

Topics

- Constraints
- Triggers
- Assertions

Constraints in SQL

- Constraints on **attribute** values:
 - these are checked whenever there is insertion to table or attribute update
 - not null constraint
 - attribute based check constraint
 - E.g., `sex char(1) CHECK (sex IN ('F', 'M'))`
 - domain constraint
 - E.g., `Create domain gender-domain CHAR (1) CHECK (VALUE IN ('F', 'M'))`
 - define sex in schema defn to be of type gender-domain

Integrity Constraints

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type. Always enforced.

Constraints in SQL

- Constraints on **tuples**
- Tuple based CHECK constraint:

```
CREATE TABLE Gamer (  
  name CHAR(30) UNIQUE  
  gender CHAR(1) CHECK (gender in ('F', 'M'))  
  age int  
  plat int  
  CHECK (age < 100 AND age > 20)  
  CHECK (plat IN (SELECT plat FROM platform))  
)
```

these are checked on insertion to relation or tuple update

Keys: Fundamental Constraint

- In the CREATE TABLE statement, use:
 - PRIMARY KEY, UNIQUE

```
CREATE TABLE Gamer (  
    name CHAR(30) PRIMARY KEY,  
    address VARCHAR(255),  
    gender CHAR(1));
```

- Or, list at end of CREATE TABLE
PRIMARY KEY (name)

Keys...

- Can use the UNIQUE keyword in same way
 - ...but for any number of attributes
 - foreign keys, which reference attributes of a second relation, only reference PRIMARY KEY
- Indexing Keys

```
CREATE UNIQUE INDEX UserIndex ON  
  Twitter_User(screen_name)
```

- Makes insertions easier to check for key constraints

Referential Integrity Constraints

- 2 rules for Foreign Keys:

Movies(MovieName, year)

ActedIn(ActorName, MovieName)

- 1) Foreign Key must be a reference to a valid value in the referenced table.
- 2) ... must be a PRIMARY KEY in the referenced table.

Declaring FK Constraints

- **FOREIGN KEY** <attributes> **REFERENCES** <table>
(<attributes>)

```
CREATE TABLE ActedIn (  
    Name CHAR(30) PRIMARY KEY,  
    MovieName CHAR(30)  
        REFERENCES Movies(MovieName));
```

- Or, summarize at end of CREATE TABLE

```
FOREIGN KEY MovieName REFERENCES Movies(MovieName)
```

- MovieName must be a PRIMARY KEY

Declaring FK Constraints

- **FOREIGN KEY** <attributes> **REFERENCES** <table>
(<attributes>)

```
CREATE TABLE ActedIn (  
    Name CHAR(30) PRIMARY KEY,  
    MovieName CHAR(30)  
        REFERENCES Movies(MovieName));
```

- Or, summarize at end of CREATE TABLE

```
FOREIGN KEY MovieName REFERENCES Movies(MovieName)
```

- MovieName must be a PRIMARY KEY

Declaring FK Constraints

- `CREATE TABLE IF NOT EXISTS `Exam1_Twitter_Tweets` (
 `tweet_id` bigint(20) unsigned NOT NULL
 AUTO_INCREMENT, `from_user_id` bigint(20) unsigned
 NOT NULL DEFAULT '0', `tweet` varchar(255) DEFAULT
 NULL, `geo` varchar(255) NOT NULL, `created_at`
 datetime NOT NULL DEFAULT '0000-00-00 00:00:00',
 PRIMARY KEY (`tweet_id`), KEY `tweet` (`tweet`))
 ENGINE=MyISAM DEFAULT CHARSET=utf8
 AUTO_INCREMENT=29 ;`
- `ALTER TABLE Exam1_Twitter_Tweets ADD CONSTRAINT
 Tweets_User_Id FOREIGN KEY (from_user_id) REFERENCES
 Exam1_Twitter_Users (user_id);`

Constraining Attribute Values

Constrain invalid values

NOT NULL

gender CHAR(1)

CHECK (gender IN ('F', 'M'))

GameName CHAR(30)

CHECK (GameName IN

(SELECT GameName FROM Games))

Last one not the same as REFERENCE

The check is invisible to the Games table!

Constraining Values with User Defined 'Types'

- Can define new domains to use as the attribute type...

```
CREATE DOMAIN GenderDomain CHAR(1)
```

```
CHECK (VALUE IN ('F', 'M'));
```

- Then update our attribute definition...

```
gender GenderDomain
```

More Complex Constraints...

- ...Among several attributes in one table
 - Specify at the end of CREATE TABLE

CHECK (gender = 'F' OR name NOT LIKE 'Ms.%')

Giving Names to Constraints

Why give names? In order to be able to alter constraints.

Add the keyword **CONSTRAINT** and then a name:

```
ssn CHAR(50) CONSTRAINT ssnIsKey PRIMARY KEY
```

```
CREATE DOMAIN ssnDomain INT  
CONSTRAINT ninedigits CHECK (VALUE >= 100000000  
                             AND VALUE <= 999999999)
```

```
CONSTRAINT rightage  
CHECK (age >= 0 OR status = "dead")
```

Altering Constraints

```
ALTER TABLE Product DROP CONSTRAINT positivePrice
```

```
ALTER TABLE Product ADD CONSTRAINT  
positivePrice CHECK (price >= 0)
```

```
ALTER DOMAIN ssn ADD CONSTRAINT no-leading-1s  
CHECK (value >= 2000000000)
```

```
DROP ASSERTION assert1.
```


Integrity Constraints (Review)

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type. Always enforced.
 - *Primary key and foreign key constraints*: you know them.

Triggers (Active database)

- **Trigger**: A procedure that starts automatically if specified changes occur to the DBMS
- Analog to a "daemon" that **monitors** a database for certain events to occur
- Three parts:
 - **Event** (activates the trigger)
 - **Condition** (tests whether the triggers should run)
[Optional]
 - **Action** (what happens if the trigger runs)
- Semantics:
 - When event occurs, and condition is satisfied, the action is performed.

Triggers – Event,Condition,Action

- Events could be :

BEFORE|AFTER INSERT|UPDATE|DELETE ON <tableName>

e.g.: `BEFORE INSERT ON Professor`

- Condition is SQL expression or even an SQL query (query with non-empty result means TRUE)
- Action can be many different choices :
 - SQL statements , body of PSM, and even DDL and transaction-oriented statements like “commit”.

Trigger Syntax

```
CREATE TRIGGER <triggerName>
BEFORE|AFTER    INSERT|DELETE|UPDATE
  [OF <columnList>] ON <tableName>|<viewName>
  [REFERENCING [OLD AS <oldName>] [NEW AS <newName>]]
  [FOR EACH ROW] (default is "FOR EACH STATEMENT")
  [WHEN (<condition>)]
<PSM body>;
```

Example Trigger

Assume our DB has a relation schema :

Professor (pNum, pName, salary)

We want to write a trigger that :

Ensures that any new professor inserted has salary \leq 60000

Example Trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor
```

```
FOR EACH ROW
```

```
BEGIN
```

```
IF (:new.salary >= 60000)
```

```
    THEN RAISE_APPLICATION_ERROR (-20004,  
    Professor Salary');
```

'Violation of Minimum

```
END IF;
```

```
END;
```

Example trigger

```
CREATE TRIGGER minSalary BEFORE INSERT ON Professor      FOR
    EACH ROW

DECLARE temp int;  -- dummy variable not needed

BEGIN
    IF (:new.salary >= 60000)
        THEN RAISE APPLICATION_ERROR (-20004,          'Violation of
Minimum Professor Salary');
    END IF;

temp := 10;      -- to illustrate declared variables

END;

.
```

Details of Trigger Example

- BEFORE INSERT ON Professor
 - This trigger is checked before the tuple is inserted
- FOR EACH ROW
 - specifies that trigger is performed for each row inserted
- :new
 - refers to the new tuple inserted
- If (:new.salary >= 60000)
 - then an application error is raised and hence the row is not inserted; otherwise the row is inserted.
- Use error code: -20004;
 - this is in the valid range

Row vs Statement Level Trigger

- **Row** level: activated once per modified tuple
- **Statement** level: activate once per SQL statement
- **Row** level triggers can access new data, statement level triggers cannot always do that (depends on DBMS).
- **Statement** level triggers will be more efficient if we do not need to make row-specific decisions

Row vs Statement Level Trigger

- Example: Consider a relation schema

Account (num, amount)

where we will allow creation of new accounts only during normal business hours.

Example: Statement level trigger

```
CREATE TRIGGER MYTRIG1
BEFORE INSERT ON Account
FOR EACH STATEMENT          --- is default
BEGIN
    IF (TO_CHAR(SYSDATE,'dy') IN ('sat','sun'))
    OR
    (TO_CHAR(SYSDATE,'hh24:mi') NOT BETWEEN '08:00' AND '17:00')
    THEN
        RAISE_APPLICATION_ERROR(-20500,'Cannot create new account now
!!');
    END IF;
END;
```

When to use BEFORE/AFTER

- Based on efficiency considerations or semantics.
- Suppose we perform statement-level **after insert**, then all the rows are inserted first, then if the condition fails, and all the inserted rows must be “rolled back”
- Not very efficient !!

Summary : Trigger Syntax

```
CREATE TRIGGER <triggerName>
BEFORE|AFTER      INSERT|DELETE|UPDATE
  [OF <columnList>] ON <tableName>|<viewName>
  [REFERENCING [OLD AS <oldName>] [NEW AS <newName>]]
[FOR EACH ROW] (default is "FOR EACH STATEMENT")
[WHEN (<condition>)]
<PSM body>;
```

Constraints versus Triggers

- **Constraints** are useful for database consistency
 - Use IC when sufficient
 - More opportunity for optimization
 - Not restricted into insert/delete/update
- **Triggers** are flexible and powerful
 - Alerters
 - Event logging for auditing
 - Security enforcement
 - Analysis of table accesses (statistics)
 - Workflow and business intelligence ...
- But can be hard to understand
 - Several triggers (Arbitrary order □ unpredictable !?)
 - Chain triggers (When to stop ?)
 - Recursive triggers (Termination?)

Triggers

Enable the database programmer to specify:

- when to check a constraint,
- what exactly to do.

A trigger has 3 parts:

- An **event** (e.g., update to an attribute)
- A **condition** (e.g., a query to check)
- An **action** (deletion, update, insertion)

When the **event** happens, the system will check the **constraint**, and if satisfied, will perform the **action**.

NOTE: triggers may cause cascading effects.

Database vendors did not wait for standards with triggers!

Elements of Triggers

- Timing of action execution: before, after or instead of triggering event
- The action can refer to both the old and new state of the database.
- Update events may specify a particular column or set of columns.
- A condition is specified with a WHEN clause.
- The action can be performed either for
 - once for every tuple, or
 - once for all the tuples that are changed by the database operation.

Assertions

- Assertions are constraints over a table as a whole or multiple tables.
- General form:
`CREATE ASSERTION <name> CHECK <cond>`
- An assertion must always be true at transaction boundaries. Any modification that causes it to become false is rejected.
- Similar to tables, assertions can be dropped by a DROP command.

Example Assertion

- CREATE ASSERTION RichProf CHECK
 (NOT EXISTS
 (SELECT *
 FROM dept, emp
 WHERE emp.name = dept.mgrname AND
 emp.salary < 50000))
- This assertion correctly guarantees that each professor makes more than 50000.
- If someone made a professor whose salary is less than 50K that insertion/update to dept table will be rejected.

Declaring Assertions

- `CREATE ASSERTION` <name> `CHECK` (<condition>)

```
CREATE ASSERTION RichPres CHECK  
(NOT EXISTS  
  (SELECT *  
    FROM Studio, MovieExec  
    WHERE presC# = cert#  
      AND netWorth < 100000000))
```